




(a) Selection vs. grounding**(b) Where grounding fails****(c) Each argument costs**

■ transcript input
 ■ audio input
 ▨ adapted (dashed in c); solid = untuned base
 - - - selection chance

v5 test split: 1,581 tool calls, one per dialogue. Every prompt lists 3 candidate tools and the gold tool is always among them, so selection chance is .33. Base = untuned Gemma-3n-E4B, best of three prompting conditions. Adapted = exp4 SFT, t4 (transcript) and t6 (all-audio), full-dialogue context.