




v5 test: 1,581 dialogues, 20,422 turn-level decisions; 3,162 have an intervention target.

$\mathcal{N}_{\text{ready}}$ = 6,722 non-intervention turns already holding every argument of the reference call; the other 10,538 are the comparison set.

Base = untuned Gemma-3n-E4B and Qwen2.5-Omni-7B; adapted = exp4 SFT (t4, t6), full-dialogue context. Unparseable counts as non-intervention.